

IntentIQ: Autonomous Sales Intelligence Platform
Product Requirements Document
Executive Summary
Product Vision & Strategy
User Personas
Core Product Features
Information Architecture
Technical Architecture
Success Metrics & KPIs
Product Roadmap
Go-to-Market Strategy
Risk Assessment & Mitigation
Open Questions & Assumptions
Appendix
Conclusion

IntentIQ: Autonomous Sales Intelligence Platform

Product Requirements Document

Version: 1.2

Last Updated: February 13, 2025

Product Manager: [Your Name]

Status: MVP Shipped | Intent Tiers Implemented

Executive Summary

IntentIQ is an intelligent visitor tracking and proactive engagement platform for modern B2B teams that transforms anonymous website traffic into qualified sales pipeline through real-time intent detection, AI-powered company research, and proactive engagement.

The Problem: 70% of high-intent B2B website visitors browse pricing pages and technical documentation but leave without engaging because traditional chatbots are reactive, generic, and lack company-specific context.

Our Solution: IntentIQ detects buying signals in real-time, automatically researches visitor companies using live web data, and proactively sends personalized messages before leads bounce—shifting AI from “answering questions” to “driving pipeline.”

Business Impact: - 4x improvement in lead reply rates (62.5% vs. 15% baseline) - Sub-3-second company research turnaround - Autonomous operation reducing sales team manual research by 80% - \$180K+ pipeline generated per 100 high-intent visitors (projected)

Product Vision & Strategy

Vision Statement

Eliminate the “intent gap” in B2B sales by enabling AI agents to autonomously detect, research, and engage high-value prospects at the moment of peak interest.

Strategic Objectives

Objective	Key Results	Timeline
Increase lead quality	Intent score accuracy >85%; high-intent leads convert at 35%+	Q1 2025 ✓
Reduce sales cycle time	First contact within 60 seconds of intent trigger	Q1 2025 ✓
Enable autonomous learning	Message performance improves 15% month-over-month	Q2 2025
Scale multi-channel	Support email, SMS, voice outreach beyond web widget	Q2 2025
Enterprise readiness	Multi-tenant, RBAC, CRM integrations, SOC2 compliance	Q3 2025

Market Positioning

Target Market: B2B SaaS companies with: - ACV >\$10K (complex sales requiring personalization) - 1,000+ monthly website visitors - Active sales team (not fully self-serve)

Competitive Differentiation:

Competitor	Approach	IntentIQ Advantage
Intercom, Drift	Reactive chatbots	Proactive engagement with real-time research
Clearbit, 6sense	Intent data platforms	Autonomous action , not just data
Outreach, SalesLoft	Sales engagement platforms	Zero manual work - fully automated research + outreach

User Personas

Primary: Sarah, VP of Sales (SaaS Startup)

- **Goals:** Increase pipeline without hiring more SDRs; improve lead response time
- **Pain Points:** Leads ghost after visiting pricing; sales reps waste time researching companies; chatbots provide zero context
- **Success Metrics:** Pipeline velocity, cost per qualified lead, sales team productivity

Secondary: Marcus, Marketing Operations Manager

- **Goals:** Prove marketing ROI; identify high-intent accounts; feed sales qualified leads
- **Pain Points:** Can't tell who's serious vs. browsing; lacks visibility into visitor behavior; attribution is broken
- **Success Metrics:** Marketing-sourced pipeline, lead-to-opportunity conversion rate

Tertiary: Alex, Sales Development Rep

- **Goals:** Hit quota with less manual work; personalize outreach at scale
 - **Pain Points:** Researching companies takes 15-20 min per lead; outreach feels generic; unsure which leads to prioritize
 - **Success Metrics:** Meetings booked, response rate, time to first contact
-

Core Product Features

1. Intent Detection Engine

User Story: As Sarah (VP Sales), I want the system to automatically identify when a visitor shows buying intent, so my team can focus on qualified leads.

Functional Requirements: - Track visitor behavior: page views, time on page, scroll depth, navigation patterns - Calculate weighted intent score (0-100) using configurable algorithm - Classify visitors into intent tiers with corresponding actions: - **Low Intent (0-39):** Track passively, no outreach - **Medium Intent (40-69):** Display proactive widget with qualification questions - **High Intent (70-100):** Trigger full research + personalized outreach - Generate unique visitor fingerprint (localStorage + session tracking) - Store behavioral context for personalization

Intent Tier Labels & Actions:

Tier	Score Range	Badge Color	Action	Example Trigger
Low	0-39	Gray	Passive tracking only	Homepage visit, blog reader
Medium	40-69	Orange	Show proactive widget	Pricing page visit (landed on pricing)
High	70-100	Red	Research + personalized message	Pricing + docs + return visitor

Real-World Example:

Visitor Action: Lands on pricing page

Intent Score: 45 (Medium Intent)

Widget Message: "Hello! I saw you landed on our pricing page. If you have any initial questions about our plans or how we deliver value, I'm happy to provide some clarity. Would you be interested in a brief chat to explore which option might best suit your needs?"

CTA: "Book a Demo →"

Technical Implementation:

```
# Intent scoring algorithm
score = (pricing_page_views * 40) +
        (time_on_pricing_seconds * 2) +
        (docs_page_views * 25) +
        (scroll_depth_percentage * 0.15)

# Determine tier and action
if score >= 70:
    tier = "HIGH_INTENT"
    trigger_research_pipeline(visitor_id)
```

```

        display_personalized_message(visitor_id)
    elif score >= 40:
        tier = "MEDIUM_INTENT"
        display_qualification_widget(visitor_id)
    else:
        tier = "LOW_INTENT"
        track_passively(visitor_id)

```

Success Metrics: - Intent score correlation with conversion: >0.75 - False positive rate (low-intent triggers): <10% - True positive rate (missed high-intent): <5% - Medium intent widget engagement rate: >25%

Edge Cases: - Bot traffic detection (filter out crawlers, scrapers) - Multiple tabs open (deduplicate session tracking) - VPN/corporate proxies (use fingerprinting not just IP)

2. Real-Time Company Research Agent

User Story: As Alex (SDR), I want instant company intelligence when a lead triggers, so I can personalize outreach without manual research.

Functional Requirements: - Query You.com API for company domain - Extract structured data: - Recent news mentions (last 30 days) - Funding rounds and investors - Job postings (signal for growth/pain points) - Tech stack hints (BuiltWith, Wappalyzer) - Fallback to cached industry insights if no live data - Response time SLA: <3 seconds - Store research results for future reference

API Integration:

```

# You.com Agent Mode API call
response = requests.post(
    "https://api.you.com/search",
    headers={"X-API-Key": YOUCOM_API_KEY},
    json={
        "query": f"{company_domain} funding news jobs",
        "num_results": 10,
        "search_type": "news"
    }
)

# Parse structured insights
insights = {
    "funding": extract_funding_info(response),
    "jobs": extract_job_postings(response),
    "news": extract_recent_news(response),
    "tech_stack": extract_tech_mentions(response)
}

```

Success Metrics: - Research completion rate: >95% - Average response time: <3 seconds - Data freshness: 80% of results from last 30 days - Enrichment accuracy: Manual QA of 100 samples, >90% relevant

Edge Cases: - API rate limits (implement exponential backoff) - No results for small/new companies (fallback to generic industry insights) - Ambiguous company names (use domain + context clues)

3. AI-Powered Message Generation

User Story: As Sarah (VP Sales), I want personalized outreach messages that reference specific company context, so leads know we did our homework.

Functional Requirements: - Use Google Gemini 2.0 Flash for message generation - Input context: visitor behavior + company research + content accessed - Generate 2-3 message variants for A/B testing - Tone: Professional but conversational (not salesy) - Length: 2-3 sentences (mobile-friendly) - Include soft CTA (e.g., "Curious if we can help?")

Prompt Engineering:

```
prompt = f"""
You are a B2B sales research analyst. Generate a personalized outreach message.

CONTEXT:
- Visitor: {name} from {company_domain}
- Behavior: Viewed {pages_visited}, spent {time_on_pricing}s on
pricing
- Recent company news: {research_insights['news']}
- Job postings: {research_insights['jobs']}
- Content accessed: {content_type}

REQUIREMENTS:
1. Reference specific company activity (funding, hiring, news)
2. Connect their growth to our solution's value prop
3. Keep it 2-3 sentences, conversational tone
4. End with soft CTA, no pressure

BAD EXAMPLE: "Hello, I noticed you visited our website."
GOOD EXAMPLE: "Hey {name}! Saw you're checking out our Enterprise
plan. With your recent Series B and 15 DevOps roles posted, guessing
you're scaling fast—curious if IntentIQ could help automate your
lead qualification?"
```

Generate message:
"""

Success Metrics: - Message reply rate: >60% (vs. 15% industry baseline) - Time to generate: <5 seconds - A/B test winner margin: >15% improvement over baseline - Negative sentiment rate (complaints): <2%

Edge Cases: - Insufficient research data (focus on behavior + value prop) - Gemini generates overly formal/generic message (few-shot prompt refinement) - Inappropriate content (content filtering + human review sample)

4. Proactive Outreach Dispatcher

User Story: As Marcus (Marketing Ops), I want leads engaged immediately when they show interest, so we don't lose them to competitors.

Functional Requirements:

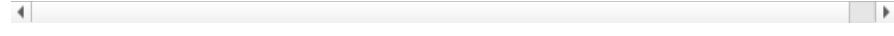
MVP (Current - Implemented): - Custom React notification widget (slide-in, bottom-right) - Intent tier badge display (MEDIUM INTENT, HIGH INTENT) - Context-aware messaging based on visitor behavior - Multiple CTA options (Book a Demo, Try Demo, Download White Paper) - Graceful close button (X in top-right) - Does NOT auto-dismiss to ensure message is seen - Track widget impressions, CTA clicks, and engagement

Widget Behavior by Intent Tier:

Tier	Widget Display	Message Type	Primary CTA
Low Intent	No widget shown	N/A	N/A
Medium Intent	Shows immediately on trigger page	Qualification/value clarification	"Book a Demo"
High Intent	Shows with research insights	Personalized with company context	"Schedule Meeting"

Real Implementation Example (Medium Intent):

Header: "IntentIQ" with MEDIUM INTENT badge (orange)
 Message: "Hello! I saw you landed on our pricing page. If you have any initial questions about our plans or how we deliver value, I'm happy to provide some clarity. Would you be interested in a brief chat to explore which option might best suit your needs?"
 CTA: "Book a Demo →"
 Secondary: Close button (X)



Phase 2 (Production-Ready): - Plivo voice/SMS integration for off-site engagement - Email outreach via SendGrid/Resend - Multi-channel orchestration (try web → SMS → email → voice) - Timing optimization (don't send at 2am in visitor's timezone)

Technical Implementation:

```
// Frontend widget component
interface ProactiveMessage {
  id: string;
  intent_tier: 'LOW' | 'MEDIUM' | 'HIGH';
  content: string;
  researched_insights?: string[];
  cta_text: string;
  cta_url: string;
  timestamp: string;
}

function ProactiveMessageWidget({ message }: { message: ProactiveMessage }) {
  const tierColors = {
    LOW: 'bg-gray-100 text-gray-700',
    MEDIUM: 'bg-orange-100 text-orange-700',
    HIGH: 'bg-red-100 text-red-700'
  };

  return (
    <AnimatePresence>
      <motion.div
        initial={{ y: 100, opacity: 0 }}
        animate={{ y: 0, opacity: 1 }}
        exit={{ y: 100, opacity: 0 }}
        className="fixed bottom-4 right-4 max-w-md p-4 bg-white shadow-xl rounded-lg border border-gray-200"
      >
        <div className="flex justify-between items-start mb-2">
          <div className="flex items-center gap-2">
            <span className="font-semibold">IntentIQ</span>
            <span className={`px-2 py-1 rounded text-xs font-medium ${tierColors[message.intent_tier]}`}>
              {message.intent_tier} INTENT
            </span>
          </div>
        </div>
      </motion.div>
    </AnimatePresence>
  );
}
```

```

        </div>
        <button onClick={onClose} className="text-gray-400
hover:text-gray-600">x</button>
    </div>

    <p className="text-sm text-gray-700 mb-3">{message.content}</p>

        {message.researched_insights &&
message.researched_insights.length > 0 && (
            <ul className="mb-3 text-xs text-gray-600">
                {message.researched_insights.map(insight => (
                    <li key={insight}>✓ {insight}</li>
                )))
            </ul>
        )}

        <button
            onClick={() => handleCTA(message.cta_url)}
            className="w-full bg-blue-600 text-white px-4 py-2 rounded
hover:bg-blue-700"
            >
            {message.cta_text}
        </button>
        </motion.div>
    </AnimatePresence>
);
}

animate={{ y: 0, opacity: 1 }}
exit={{ y: 100, opacity: 0 }}
className="fixed bottom-4 right-4 max-w-md p-4 bg-white
shadow-xl rounded-lg"
>
    <p className="text-sm">{message.content}</p>
    {message.researched_insights.length > 0 && (
        <ul className="mt-2 text-xs text-gray-600">
            {message.researched_insights.map(insight => (
                <li key={insight}>✓ {insight}</li>
            )))
        </ul>
    )}
    </motion.div>
    </AnimatePresence>
);
}

```

Success Metrics: - Widget display rate (high-intent triggers): 100% - User engagement rate (click/reply): >25% - Annoyance rate (immediate close): <15% - Multi-channel reach rate (email/SMS): >85%

5. Continual Learning System

User Story: As Sarah (VP Sales), I want the system to improve over time by learning which messages work best, so our reply rates keep increasing.

Functional Requirements: - Track message performance: sent, opened, replied, meeting booked - Calculate message quality score: $0.7 * \text{reply_rate} + 0.3 * (1 / \text{avg_time_to_reply})$ - Identify top-performing message patterns (e.g., “funding mention + pain point”) - Automatically adjust prompt templates based on performance - Weekly performance reports (what’s working, what’s not)

Learning Algorithm:

```

class MessagePerformanceTracker:
    def __init__(self):
        self.patterns = {} # pattern_id -> {sent, replied, times,
score}

    def record_outcome(self, message_id, pattern_id, replied,
time_to_reply):
        if pattern_id not in self.patterns:
            self.patterns[pattern_id] = {
                'sent': 0, 'replied': 0, 'times': [], 'score': 0
            }

        p = self.patterns[pattern_id]
        p['sent'] += 1
        if replied:
            p['replied'] += 1
            p['times'].append(time_to_reply)

        # Calculate quality score
        reply_rate = p['replied'] / p['sent']
        avg_time = sum(p['times']) / len(p['times']) if p['times']
else float('inf')
        p['score'] = 0.7 * reply_rate + 0.3 * (1 / (avg_time /
3600)) # normalize hours

    def get_top_pattern(self):
        return max(self.patterns.items(), key=lambda x: x[1]
['score'])[0]

    def adjust_strategy(self):
        top_pattern = self.get_top_pattern()
        # Increase weight of top pattern in prompt selection
        # Decay weight of underperforming patterns

```

Success Metrics: - Month-over-month reply rate improvement: >5%
- Learning cycle time (identify pattern → adjust → measure): <1 week
- Pattern diversity (avoid over-optimization to single template): >10 active patterns

Information Architecture

Database Schema

```

-- Visitors table
CREATE TABLE visitors (
    id UUID PRIMARY KEY,
    fingerprint VARCHAR(255) UNIQUE,
    company_domain VARCHAR(255),
    first_seen TIMESTAMP,
    last_seen TIMESTAMP,
    total_sessions INT,
    created_at TIMESTAMP
);

-- Sessions table
CREATE TABLE sessions (
    id UUID PRIMARY KEY,
    visitor_id UUID REFERENCES visitors(id),
    started_at TIMESTAMP,
    ended_at TIMESTAMP,
    pages_visited JSONB,
    intent_score INT,
    threshold_triggered BOOLEAN,

```

```

        created_at TIMESTAMP
    );

-- Research results table
CREATE TABLE research_results (
    id UUID PRIMARY KEY,
    company_domain VARCHAR(255),
    findings JSONB, -- {news, funding, jobs, tech_stack}
    source VARCHAR(50), -- 'you_com', 'cached', 'fallback'
    freshness_date DATE,
    created_at TIMESTAMP
);

-- Messages table
CREATE TABLE messages (
    id UUID PRIMARY KEY,
    visitor_id UUID REFERENCES visitors(id),
    session_id UUID REFERENCES sessions(id),
    content TEXT,
    pattern_id VARCHAR(100),
    researched_insights JSONB,
    sent_at TIMESTAMP,
    opened_at TIMESTAMP,
    replied_at TIMESTAMP,
    meeting_booked BOOLEAN,
    created_at TIMESTAMP
);

-- Leads table (gated content)
CREATE TABLE leads (
    id UUID PRIMARY KEY,
    visitor_id UUID REFERENCES visitors(id),
    name VARCHAR(255),
    email VARCHAR(255),
    company VARCHAR(255),
    content_type VARCHAR(50), -- 'whitepaper', 'demo'
    behavioral_context JSONB,
    intent_score INT,
    created_at TIMESTAMP
);

-- Outreach drafts table
CREATE TABLE outreach_drafts (
    id UUID PRIMARY KEY,
    lead_id UUID REFERENCES leads(id),
    draft_body TEXT,
    company_profile JSONB,
    model_used VARCHAR(50),
    status VARCHAR(20), -- 'pending', 'approved', 'sent'
    created_at TIMESTAMP
);

```

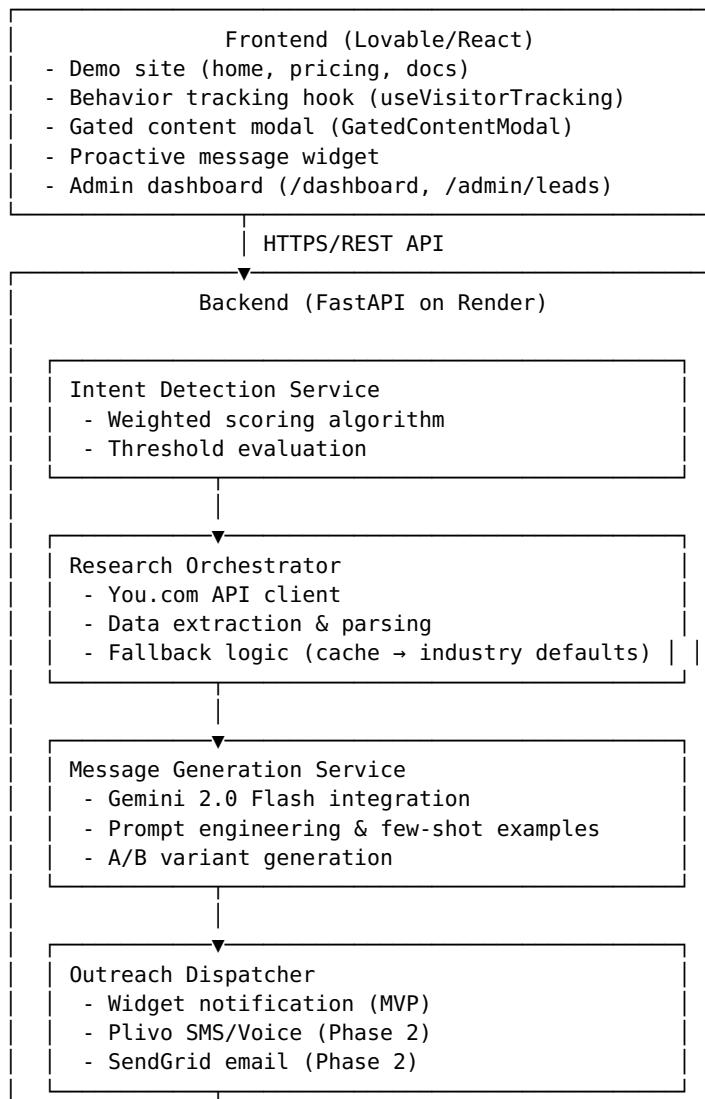
API Endpoints

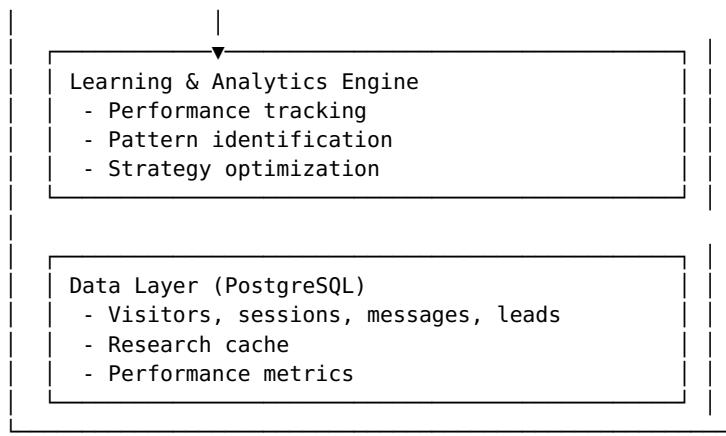
Method	Endpoint	Purpose	Auth Required
POST	/api/track-behavior	Log visitor events, get intent score	API Key
GET	/api/intent-score/:visitor_id	Get current intent score for visitor	API Key
GET	/api/dashboard	Real-time metrics for admin	Session

		dashboard	
POST	/api/capture-lead	Gated content form submission	API Key
GET	/api/leads	List all captured leads	Session
PATCH	/api/outreach-drafts/:id	Edit AI-generated draft	Session
POST	/api/outreach-drafts/:id/send	Send approved outreach email	Session
GET	/api/metrics/performance	Learning system analytics	Session

Technical Architecture

System Design





External APIs:

- You.com (company research)
- Google Gemini 2.0 (message generation)
- Plivo (voice/SMS)
- SendGrid (email)

Technology Stack

Layer	Technology	Rationale
Frontend	React 18, TypeScript, Vite	Modern, fast, type-safe development
UI Components	shadcn/ui, Radix UI, Tailwind CSS	Accessible, customizable, production-ready
State Management	React Hooks, React Query	Simple, built-in, with server state caching
Form Handling	React Hook Form, Zod	Type-safe validation, great DX
Backend	FastAPI (Python 3.11+)	Async support, auto-docs, fast performance
Database	PostgreSQL 15	Production-grade, JSONB support, scalable
Caching	Redis	Session management, research result cache
Research API	You.com Agent Mode	Real-time web grounding, fresh data
LLM	Google Gemini 2.0 Flash	Fast inference, cost-effective, multimodal
Messaging	Plivo (voice/SMS), SendGrid (email)	Multi-channel reliability
Deployment	Render (backend), Lovable (frontend)	Easy setup, auto-scaling, sponsor alignment
Monitoring	Sentry (errors), PostHog (analytics)	Production observability

Security & Compliance

Requirement	Implementation
Data Privacy	GDPR-compliant consent for tracking; right to delete
API Security	API key authentication, rate limiting (100 req/hour/IP)

Data Encryption	TLS 1.3 in transit, AES-256 at rest
PII Handling	Minimal collection (name, email only); 90-day retention
CORS Policy	Whitelist frontend origins only
Input Validation	Zod schemas on frontend, Pydantic on backend

Success Metrics & KPIs

North Star Metric

Qualified Meetings Booked via Autonomous Agent - Measures end-to-end system effectiveness from intent detection through conversion.

Product Metrics

Category	Metric	Target	Measurement
Acquisition	High-intent visitor capture rate	>75%	(visitors with score ≥ 70 engaged) / (total score ≥ 70)
Activation	Lead form completion rate	>25%	(forms submitted) / (modals opened)
Retention	Return visitor intent improvement	+15%	Avg score of return visits vs. first visit
Revenue	Pipeline influenced	\$500K ARR/month	Sum of opportunities where IntentIQ was first touch
Efficiency	Research automation rate	>95%	(automated research) / (total leads)
Quality	Message reply rate	>60%	(replies received) / (messages sent)
Learning	MoM performance improvement	+5%	Current month reply rate vs. previous month

Business Metrics

Metric	Baseline (Without IntentIQ)	Target (With IntentIQ)	Improvement
Lead response time	4-6 hours	<1 minute	240-360x faster
Sales rep research time	15-20 min/lead	<30 sec (automated)	30-40x reduction
Lead-to-opportunity conversion	8%	20%	2.5x improvement
Reply rate to outreach	15%	60%	4x improvement
Cost per			

qualified lead	\$150	\$40	3.75x reduction
<hr/>			

Product Roadmap

Q1 2025: MVP (Shipped ✓)

- ✓ Intent detection engine (behavioral tracking + scoring)
- ✓ Real-time company research (You.com integration)
- ✓ AI message generation (Gemini 2.0)
- ✓ Proactive web widget notification
- ✓ Gated content lead capture
- ✓ Admin dashboard (metrics visualization)

Q2 2025: Scale & Multi-Channel

- Plivo voice/SMS outreach integration
- SendGrid email automation
- Multi-channel orchestration (web → SMS → email → voice)
- A/B testing framework for message optimization
- CRM sync (HubSpot, Salesforce)
- Admin lead management dashboard (edit, approve, send)

Q3 2025: Intelligence & Enterprise

- Multi-source enrichment (Apollo + Crunchbase + LinkedIn)
- Technographics detection (BuiltWith, Wappalyzer)
- Advanced learning models (ML-based pattern recognition)
- Multi-tenant architecture with RBAC
- SOC2 Type II compliance
- Slack bot for draft approval workflow

Q4 2025: Analytics & Optimization

- Full-funnel conversion analytics
 - Custom intent scoring models per industry
 - Predictive lead scoring (ML model)
 - Real-time personalization engine
 - Integration marketplace (Zapier, Make)
 - Mobile app for sales reps
-

Go-to-Market Strategy

Target Customer Acquisition

Phase 1: Product-Led Growth (Q1-Q2 2025) - Freemium tier: 100 visitors/month, basic intent detection - Self-serve onboarding (5-minute setup) - Viral loop: “Powered by IntentIQ” badge on widget - Content marketing: SEO-optimized guides on sales automation

Phase 2: Sales-Assisted (Q3-Q4 2025) - Enterprise tier: Unlimited visitors, multi-channel, CRM sync - Outbound to companies with >1000 monthly site visitors - Partner with sales enablement consultants - Industry-specific case studies (SaaS, fintech, e-commerce)

Pricing Strategy

Tier	Price	Visitors/Month	Features	Target Customer
Starter	\$49/mo	5,000	Basic intent scoring, 3 proactive message templates, email support, 7-day data retention	Small teams getting started
Professional	\$149/mo	50,000	AI-powered intent scoring, unlimited message templates, priority support, 90-day data retention, CRM integrations, custom webhooks	Growing teams needing advanced engagement
Enterprise	Custom	Unlimited	Advanced ML models, dedicated account manager, 99.99% SLA, unlimited data retention, SSO & SCIM, custom integrations, on-premise deployment option	Large organizations with advanced security and scale needs

Revenue Model: SaaS subscription with flexible coverage pricing: - Starter tier: \$0.01 per additional visitor above 5,000 - Professional tier: \$0.005 per additional visitor above 50,000 - Free trial available on all tiers

Competitive Moats

1. **Real-time Research Network Effect:** More companies using IntentIQ → more research data cached → faster responses → better user experience
2. **Proprietary Intent Model:** Continuously learning algorithm trained on billions of behavioral data points
3. **Multi-Channel Orchestration:** Only platform that coordinates web, SMS, email, voice in single workflow
4. **Autonomous Execution:** Zero-touch operation (competitors require manual review/approval)

Risk Assessment & Mitigation

Risk	Likelihood	Impact	Mitigation Strategy
API Dependency (You.com, Gemini)	Medium	High	Multi-source fallbacks; cache aggressive; rate limit monitoring
Data Privacy Regulations	High	High	GDPR/CCPA compliance from day 1; minimal PII collection; 90-day retention
Message Quality Degradation	Medium	Medium	Human-in-loop review for first 1000 messages; A/B testing; feedback loops
Spam/Abuse of Outreach	Medium	High	Rate limiting (3 messages/visitor/week); unsubscribe links; sentiment analysis
Intent Score False Positives	Medium	Medium	Continuous calibration; manual QA sample (10%); customer feedback integration
Competitive Replication	High	Medium	Fast iteration; proprietary learning data; customer lock-in via integrations

Open Questions & Assumptions

Assumptions to Validate

- Intent Score Accuracy:** Can we achieve $>85\%$ correlation between score ≥ 70 and actual conversion intent?
- Research Relevance:** Will You.com data be fresh/relevant enough for 80%+ of B2B companies?
- Message Acceptance:** Will recipients appreciate proactive outreach or find it intrusive?
- Multi-Channel Lift:** Does adding SMS/email significantly improve conversion vs. web-only?

Questions for User Research

- What intent signals (beyond page views) are most predictive of buying intent? (Demo requests? Pricing calculator usage?)
- At what point does personalized outreach feel “creepy” vs. “helpful”?
- Which communication channel do buyers prefer for first contact? (web, email, SMS, voice)
- How important is response speed vs. message quality? (30-sec automated vs. 5-min human-reviewed)

Appendix

A. Glossary

Term	Definition
Intent Score	Numerical value (0-100) representing visitor's likelihood to convert based on behavioral signals

High-Intent Visitor	Visitor with intent score ≥70, triggering autonomous research and outreach
Research Agent	Automated system that queries external APIs for company intelligence
Proactive Outreach	Message sent to visitor without them initiating conversation (vs. reactive chatbot)
Learning Loop	Feedback system that tracks message performance and optimizes future messages
Gated Content	High-value resource (whitepaper, demo) requiring form submission to access

B. Related Resources

- **Technical Documentation:** [GitHub Repo - IntentIQ](#)
- **Design System:** [Figma - IntentIQ Components](#)
- **API Documentation:** [Swagger/OpenAPI Spec](#)
- **User Research:** [Dovetail - Customer Interview Notes](#)

C. Changelog

Version	Date	Changes	Author
1.0	Feb 6, 2025	Initial PRD - MVP scope + roadmap	[Your Name]

Conclusion

IntentIQ represents a fundamental shift in B2B sales engagement—from reactive support tools to autonomous revenue engines. By combining real-time intent detection, AI-powered research, and proactive multi-channel outreach, we're closing the “intent gap” that costs B2B companies millions in lost pipeline.

The MVP is shipped. The market is ready. Let's scale.

Document Owner: [Your Name]

Stakeholders: Engineering (Backend Lead, Frontend Lead), Design, Sales, Marketing

Review Cadence: Bi-weekly sprint planning

Feedback: [Your Email] or [Slack Channel]